

ujjwal-singh-project-3

June 12, 2024

Meta Stock Price Prediction

0.0.1 Institution Name: Vigor Council

0.0.2 Guidance Under: Dr. B.P. Sharma

0.0.3 Intern Name: Ujjwal Singh

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[2]: df=pd.read_csv("META.csv")
```

```
[3]: df.head()
```

```
[3]:
```

	Date	Open	High	Low	Close	Adj Close	\
0	28/10/2021	312.989990	325.519989	308.109985	316.920013	316.584106	
1	29/10/2021	320.190002	326.000000	319.600006	323.570007	323.227051	
2	01/11/2021	326.040009	333.450012	326.000000	329.980011	329.630280	
3	02/11/2021	331.380005	334.790009	323.799988	328.079987	327.732269	
4	03/11/2021	327.489990	332.149994	323.200012	331.619995	331.268524	

	Volume
0	50806800
1	37059400
2	31518900
3	28353000
4	20786500

```
[4]: df.columns
```

```
[4]: Index(['Date', 'Open', 'High', 'Low', 'Close', 'Adj Close', 'Volume'],
dtype='object')
```

```
[5]: df.shape
```

```
[5]: (633, 7)
```

```
[6]: df.isnull().sum()
```

```
[6]: Date          0
      Open         0
      High         0
      Low          0
      Close        0
      Adj Close    0
      Volume       0
      dtype: int64
```

```
[7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 633 entries, 0 to 632
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Date        633 non-null   object
 1   Open        633 non-null   float64
 2   High        633 non-null   float64
 3   Low         633 non-null   float64
 4   Close       633 non-null   float64
 5   Adj Close   633 non-null   float64
 6   Volume      633 non-null   int64
dtypes: float64(5), int64(1), object(1)
memory usage: 34.7+ KB
```

```
[8]: df['Date']=pd.to_datetime(df["Date"], format='%d/%m/%Y')
```

```
[9]: Date_column=df.columns[0]
```

```
[10]: df.set_index(Date_column, inplace=True)
```

```
[11]: df.head()
```


```
[11]:
```

	Open	High	Low	Close	Adj Close	\
Date						
2021-10-28	312.989990	325.519989	308.109985	316.920013	316.584106	
2021-10-29	320.190002	326.000000	319.600006	323.570007	323.227051	
2021-11-01	326.040009	333.450012	326.000000	329.980011	329.630280	
2021-11-02	331.380005	334.790009	323.799988	328.079987	327.732269	
2021-11-03	327.489990	332.149994	323.200012	331.619995	331.268524	
Volume						
Date						
2021-10-28	50806800					

```
2021-10-29  37059400
2021-11-01  31518900
2021-11-02  28353000
2021-11-03  20786500
```

```
[12]: plt.figure(figsize=(12,8))
      column="Open"
      plt.plot(df.index, df[column], label=column)

      plt.xlabel('Date')
      plt.ylabel('Values')
      plt.title(f'Time Series for {column}')
      plt.legend()
      plt.show()
```




```
[13]: plt.figure(figsize=(12,8))
      column="High"
      plt.plot(df.index, df[column], label=column,color='y')

      plt.xlabel('Date')
      plt.ylabel('Values')
      plt.title(f'Time Series for {column}')
```

```
plt.legend()  
plt.show()
```



```
[14]: plt.figure(figsize=(12,8))  
      column="Low"  
      plt.plot(df.index, df[column], label=column, color='k')  
  
      plt.xlabel('Date')  
      plt.ylabel('Values')  
      plt.title(f'Time Series for {column}')  
      plt.legend()  
      plt.show()
```




```
[15]: plt.figure(figsize=(12,8))
      column="Close"
      plt.plot(df.index, df[column], label=column,color='g')

      plt.xlabel('Date')
      plt.ylabel('Values')
      plt.title(f'Time Series for {column}')
      plt.legend()
      plt.show()
```



```
[16]: plt.figure(figsize=(12,8))
column="Adj Close"
plt.plot(df.index, df[column], label=column,color='m')

plt.xlabel('Date')
plt.ylabel('Values')
plt.title(f'Time Series for {column}')
plt.legend()
plt.show()
```



```
[17]: pip install mplfinance pandas
```

Requirement already satisfied: mplfinance in c:\users\kirti\anaconda3\lib\site-packages (0.12.10b0)Note: you may need to restart the kernel to use updated packages.

Requirement already satisfied: pandas in c:\users\kirti\anaconda3\lib\site-packages (2.0.3)

Requirement already satisfied: matplotlib in c:\users\kirti\anaconda3\lib\site-packages (from mplfinance) (3.7.2)

Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\kirti\anaconda3\lib\site-packages (from pandas) (2.8.2)

Requirement already satisfied: pytz>=2020.1 in c:\users\kirti\anaconda3\lib\site-packages (from pandas) (2023.3.post1)

Requirement already satisfied: tzdata>=2022.1 in c:\users\kirti\anaconda3\lib\site-packages (from pandas) (2023.3)

Requirement already satisfied: numpy>=1.21.0 in c:\users\kirti\anaconda3\lib\site-packages (from pandas) (1.24.3)

Requirement already satisfied: six>=1.5 in c:\users\kirti\anaconda3\lib\site-packages (from python-dateutil>=2.8.2->pandas) (1.16.0)

Requirement already satisfied: contourpy>=1.0.1 in c:\users\kirti\anaconda3\lib\site-packages (from matplotlib->mplfinance) (1.0.5)

Requirement already satisfied: cycycler>=0.10 in

```
c:\users\kirti\anaconda3\lib\site-packages (from matplotlib->mplfinance)
(0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in
c:\users\kirti\anaconda3\lib\site-packages (from matplotlib->mplfinance)
(4.25.0)
Requirement already satisfied: kiwisolver>=1.0.1 in
c:\users\kirti\anaconda3\lib\site-packages (from matplotlib->mplfinance) (1.4.4)
Requirement already satisfied: packaging>=20.0 in
c:\users\kirti\anaconda3\lib\site-packages (from matplotlib->mplfinance) (23.1)
Requirement already satisfied: pillow>=6.2.0 in
c:\users\kirti\anaconda3\lib\site-packages (from matplotlib->mplfinance)
(10.2.0)
Requirement already satisfied: pyparsing<3.1,>=2.3.1 in
c:\users\kirti\anaconda3\lib\site-packages (from matplotlib->mplfinance) (3.0.9)
```

```
[18]: import mplfinance as mpf
import warnings
warnings.filterwarnings("ignore")
```

```
[19]: plt.figure(figsize=(25,6))
mpf.plot(data=df, type='candle', volume=True, style='yahoo')
```

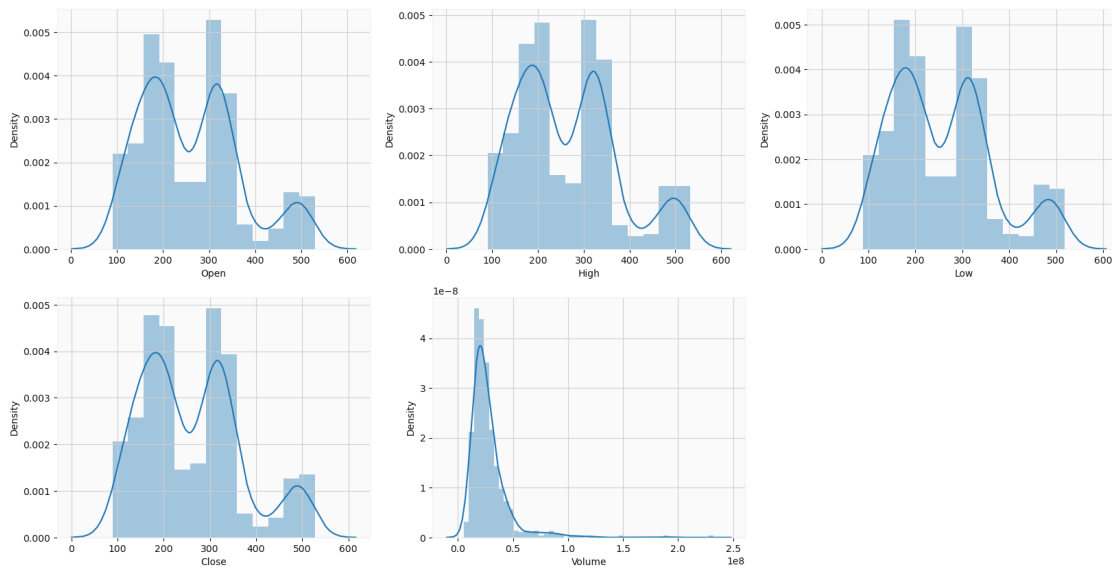
<Figure size 2500x600 with 0 Axes>



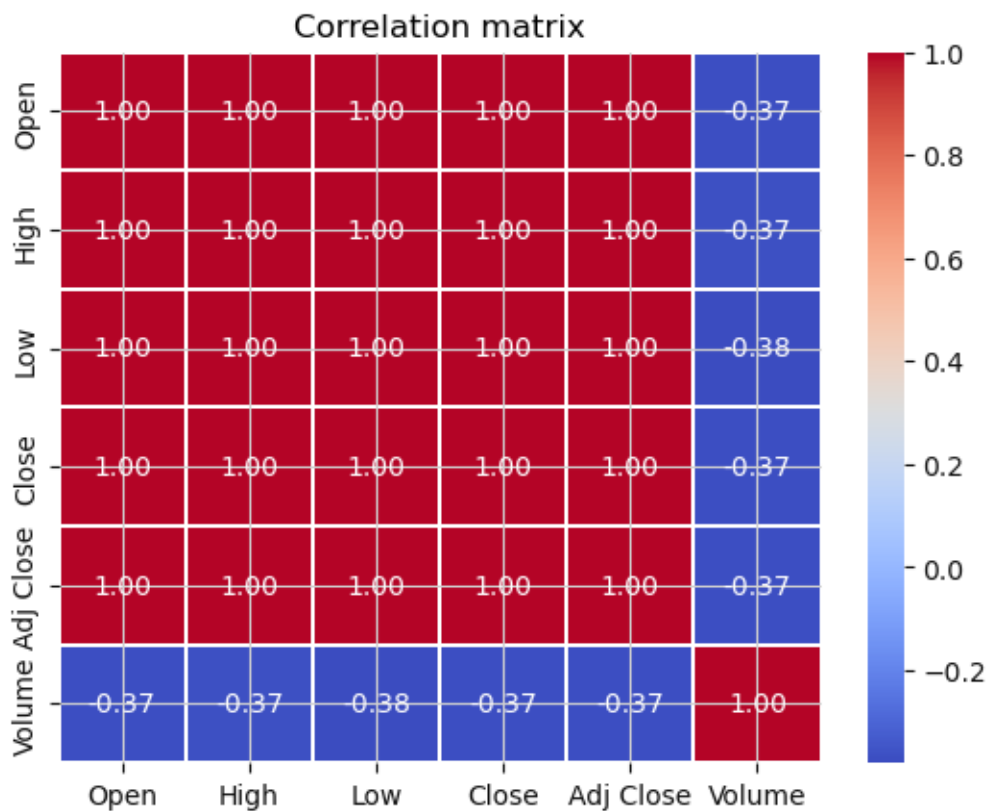

```
[20]: features = ['Open', 'High', 'Low', 'Close', 'Volume']

plt.subplots(figsize=(20,10))

for i, col in enumerate(features):
    plt.subplot(2,3,i+1)
    sns.distplot(df[col], kde=True)
plt.show()
```




```
[21]: correlation_matrix=df.corr()
sns.heatmap(correlation_matrix, cmap="coolwarm", annot=True, fmt='.2f',
            ↳linewidths=.25)
plt.title("Correlation matrix")
plt.show()
```



```
[22]: df['MA50'] = df['Close'].rolling(window=50).mean()
df['MA200'] = df['Close'].rolling(window=200).mean()

plt.figure(figsize=(10, 6))
plt.plot( df['Close'], label='Close Price')
plt.plot( df['MA50'], label='50-Day Moving Average')
plt.plot( df['MA200'], label='200-Day Moving Average')
plt.xlabel('Date')
plt.ylabel('Price')
plt.title('Meta Stock Prices with Moving Averages')
plt.legend()
plt.show()
```



```
[23]: df=df.drop(["MA50","MA200"], axis=1)
```

0.0.4 Machine Learning Models¶

These models can be applied to regression tasks where the goal is to predict a continuous target variable based on one or more input features.

- 1.Linear Regression
- 2.Ridge Regression
- 3.Lasso Regression
- 4.ElasticNet Regression
- 5.Support Vector Machines (SVM) with kernel functions like linear, polynomial, or RBF
- 6.Decision Trees (and ensemble methods like Random Forests)
- 7.Gradient Boosting Machines (GBM) and its variants like XGBoost, LightGBM, and CatBoost
- 8.Neural Networks (e.g., Multi-layer Perceptron, Convolutional Neural Networks for image data, Recurrent Neural Networks for sequential data)
- 9.K-Nearest Neighbors Regression(KNN)

0.0.5 Suitable metrics

- 1.Mean Absolute Error (MAE)
- 2.Mean Squared Error (MSE)
- 3.Root Mean Squared Error (RMSE)
- 4.Mean Absolute Percentage Error (MAPE)
- 5.R-squared (R^2)

6.Adjusted R-squared

1 Using Linear Regression

```
[24]: from sklearn.model_selection import train_test_split
      from sklearn.linear_model import LinearRegression
      from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error
```

```
[25]: X=df.drop('Adj Close', axis=1)
      y=df['Adj Close']
```

```
[26]: X_train,X_test,y_train,y_test=train_test_split(X,y, test_size=0.3,
      ↪random_state=42)
```

```
[27]: print('X Training Shape:', X_train.shape)
      print('y Training Shape:', y_train.shape)
      print('X Testing Shape:', X_test.shape)
      print('y Testing Shape:', y_test.shape)
```

```
X Training Shape: (443, 5)
y Training Shape: (443,)
X Testing Shape: (190, 5)
y Testing Shape: (190,)
```

```
[28]: model=LinearRegression()
      model
```

```
[28]: LinearRegression()
```

```
[29]: model.fit(X_train, y_train)
```

```
[29]: LinearRegression()
```

```
[30]: model.intercept_
```

```
[30]: -0.20678568149986631
```

```
[31]: model.coef_
```

```
[31]: array([ 7.64454021e-03, -3.98022156e-03, -7.23159977e-03,  1.00335347e+00,
        4.31480664e-10])
```

```
[32]: y_pred=model.predict(X_test)
      y_pred
```

```
[32]: array([171.84560411, 129.79973625, 246.59788953, 191.38828965,
        308.41809773, 328.91001157, 169.27526613, 207.48877787,
```

334.46323582, 158.53635026, 111.18130445, 334.20369297,
 317.90381691, 198.27416881, 472.98708908, 216.30167272,
 329.70112328, 88.71677445, 141.27426537, 174.70425589,
 203.55949194, 117.81781239, 331.52620501, 438.86162709,
 326.92855877, 223.07660157, 209.17526729, 113.92908556,
 429.85157322, 111.68679373, 299.89017818, 301.37886971,
 132.59866151, 454.45936071, 195.98392525, 173.18349606,
 186.27864662, 501.45569125, 175.36863825, 177.72390232,
 215.85644666, 334.97233113, 202.76472145, 294.04653037,
 172.20098893, 165.14552966, 116.87502707, 162.83142342,
 205.93866925, 200.45321227, 156.94357382, 119.54064759,
 327.47093615, 163.25207865, 505.2283313 , 483.8293056 ,
 190.99649647, 320.27118286, 233.2760649 , 329.50388767,
 501.98273836, 312.30793678, 178.11340152, 345.9310479 ,
 129.51630146, 121.36208633, 319.2957724 , 263.37323684,
 354.55425142, 333.93801224, 358.39878179, 310.14142448,
 210.25153988, 116.66463464, 169.13117323, 334.42266439,
 370.18734672, 219.97902062, 218.6242831 , 96.26771509,
 162.66853375, 299.88272396, 281.56798152, 286.79169288,
 160.12088049, 509.28436906, 511.8612973 , 182.85612852,
 157.83141034, 166.43453556, 115.68323459, 193.79322586,
 329.53589169, 191.38999548, 219.32512118, 114.26742064,
 199.56426449, 101.28151882, 132.76387623, 122.20604054,
 294.06483514, 137.28640759, 201.95659021, 346.07392199,
 133.56778356, 207.36664839, 299.96940306, 231.59954042,
 369.44678042, 140.19679666, 301.46359284, 168.7175905 ,
 227.81958765, 188.45182357, 181.472698 , 217.06914819,
 505.68365238, 242.94426155, 393.86413683, 160.15969798,
 113.01585765, 205.1126129 , 132.66563539, 149.31572412,
 170.34070206, 486.74036938, 205.77631727, 340.843856 ,
 229.61760141, 213.91864504, 323.73337958, 394.48495483,
 489.86128947, 158.93493685, 316.71099802, 147.07863959,
 471.45450437, 135.14216191, 126.72193027, 246.46992429,
 191.04846855, 129.81684215, 327.54702644, 342.677413 ,
 219.34667983, 346.84698598, 322.81065291, 340.61591894,
 318.55574504, 511.61922312, 151.27289127, 167.73498301,
 211.25314406, 210.78919234, 116.91282254, 286.529857 ,
 201.84230749, 310.41345419, 305.4781015 , 491.05543698,
 183.22404069, 180.65389869, 273.1117357 , 293.892356 ,
 199.78778565, 353.83167052, 109.64777796, 129.62192843,
 327.40300343, 336.71287368, 491.57658048, 170.16037272,
 499.43722956, 117.88381967, 203.27648183, 120.22548459,
 285.01778726, 495.27240329, 136.15381573, 90.60158394,
 473.06596818, 516.66593128, 164.47044325, 136.20852356,
 196.38039101, 220.80268754, 187.39964 , 304.55769083,
 236.883367 , 158.89085567])

```
[33]: r2=r2_score(y_pred,y_test) ## Higher values indicate better model performance,
      ↪with 1 meaning perfect prediction
      print("R2 for Linear Regression Model is", r2)
```

R² for Linear Regression Model is 0.9999989713428837

```
[34]: mse=mean_squared_error(y_pred,y_test) ## Lower values indicates better model
      ↪performance
      print("MSE for Linear Regression Model is", mse)
```

MSE for Linear Regression Model is 0.013177732400519863

```
[35]: mea=mean_absolute_error(y_test, y_pred) ## Lower value indicates better model
      ↪performance
      print("Mean Absolute Error for Linear Regression is ", mea)
```

Mean Absolute Error for Linear Regression is 0.08363125599643

2 Using XgBoost Model

```
[36]: pip install xgboost
```

Requirement already satisfied: xgboost in c:\users\kirti\anaconda3\lib\site-packages (2.0.3)
Requirement already satisfied: numpy in c:\users\kirti\anaconda3\lib\site-packages (from xgboost) (1.24.3)
Requirement already satisfied: scipy in c:\users\kirti\anaconda3\lib\site-packages (from xgboost) (1.11.1)
Note: you may need to restart the kernel to use updated packages.

```
[37]: from xgboost import XGBRegressor
```

```
[38]: model2=XGBRegressor()
      model2
```


```
[38]: XGBRegressor(base_score=None, booster=None, callbacks=None,
                  colsample_bylevel=None, colsample_bynode=None,
                  colsample_bytree=None, device=None, early_stopping_rounds=None,
                  enable_categorical=False, eval_metric=None, feature_types=None,
                  gamma=None, grow_policy=None, importance_type=None,
                  interaction_constraints=None, learning_rate=None, max_bin=None,
                  max_cat_threshold=None, max_cat_to_onehot=None,
                  max_delta_step=None, max_depth=None, max_leaves=None,
                  min_child_weight=None, missing=nan, monotone_constraints=None,
                  multi_strategy=None, n_estimators=None, n_jobs=None,
                  num_parallel_tree=None, random_state=None, ...)
```

```
[39]: model2.fit(X_train, y_train)
```

```
[39]: XGBRegressor(base_score=None, booster=None, callbacks=None,  
    colsample_bylevel=None, colsample_bynode=None,  
    colsample_bytree=None, device=None, early_stopping_rounds=None,  
    enable_categorical=False, eval_metric=None, feature_types=None,  
    gamma=None, grow_policy=None, importance_type=None,  
    interaction_constraints=None, learning_rate=None, max_bin=None,  
    max_cat_threshold=None, max_cat_to_onehot=None,  
    max_delta_step=None, max_depth=None, max_leaves=None,  
    min_child_weight=None, missing=nan, monotone_constraints=None,  
    multi_strategy=None, n_estimators=None, n_jobs=None,  
    num_parallel_tree=None, random_state=None, ...)
```

```
[40]: # Make predictions  
y_pred_train = model2.predict(X_train)  
y_pred_test = model2.predict(X_test)
```

```
[41]: # Visualize the predictions  
plt.figure(figsize=(20, 6))  
plt.plot(df.index, df['Adj Close'], label='Actual')  
plt.plot(df.index[:len(y_train)], y_pred_train, label='Train Prediction')  
plt.plot(df.index[len(y_train):], y_pred_test, label='Test Prediction')  
plt.xlabel('Year')  
plt.ylabel('Adj Close')  
plt.title('Adj Close Prediction with XGBoost', fontweight='bold')  
plt.legend()  
plt.show()
```



```
[42]: r2=r2_score(y_train, y_pred_train)  
print(r2)
```

0.9999985831646808

```
[43]: r2=r2_score(y_test, y_pred_test)
      print("R2 for XGBoost Model is", r2)
```

R² for XGBoost Model is 0.9994371565478546

```
[44]: mean_error=mean_squared_error(y_test, y_pred_test)
      print("MSE of XGBoost Model is ", mean_error)
```

MSE of XGBoost Model is 7.211095237967466

3 Using Support Vector Regressor

```
[45]: from sklearn.svm import SVR
      from sklearn.preprocessing import StandardScaler
```

```
[46]: svr=SVR(kernel='rbf')
      svr
```

[46]: SVR()

```
[47]: svr.fit(X_train, y_train)
```

[47]: SVR()

```
[48]: y_pred_train = svr.predict(X_train)
      y_pred_test = svr.predict(X_test)
```

```
[49]: train_mse=mean_squared_error(y_train, y_pred_train)
      test_mse=mean_squared_error(y_test, y_pred_test)
      train_r2=r2_score(y_train, y_pred_train)
      test_r2=r2_score(y_test, y_pred_test)

      print("Training MSE: ", train_mse)
      print("Test MSE: ", test_mse)
      print("Train R2: ", train_r2)
      print("Test R2", test_r2)
```

Training MSE: 9112.31089323682

Test MSE: 10472.741332157315

Train R²: 0.1611369057498142

Test R² 0.1825771688908605

```
[50]: y_pred=svr.predict(X_test)
      r2 = r2_score(y_test, y_pred)
      print("R2 for Support Vector Regressor Model is", r2)
```

R² for Support Vector Regressor Model is 0.1825771688908605

4 Using Random Forest Regressor model

```
[51]: from sklearn.ensemble import RandomForestRegressor
```

```
[52]: rm=RandomForestRegressor()  
rm
```

```
[52]: RandomForestRegressor()
```

```
[53]: rm.fit(X_train, y_train)
```

```
[53]: RandomForestRegressor()
```

```
[54]: y_pred=rm.predict(X_test)  
y_pred
```

```
[54]: array([173.11981739, 128.4682894 , 243.46288574, 191.49511703,  
        308.6599035 , 328.25641521, 169.46769407, 208.32056675,  
        334.60937617, 158.06678842, 111.91425679, 334.10540912,  
        318.17112346, 198.0639494 , 470.31970758, 215.93309006,  
        329.82874163,  96.26086666, 142.01512193, 174.5194339 ,  
        204.37156081, 117.05480409, 331.7681974 , 441.33302655,  
        326.51676095, 223.44982028, 209.73077462, 116.34955175,  
        437.99585937, 108.20139744, 299.58913045, 301.3657434 ,  
        133.1404324 , 460.2586171 , 195.95608711, 173.75224616,  
        186.51900234, 502.21279687, 176.77993223, 178.28803253,  
        215.76846509, 334.66321857, 204.03341833, 294.35549059,  
        173.04859391, 163.77403531, 116.57221582, 162.94161499,  
        206.15106885, 199.30853042, 156.542111 , 119.07006527,  
        327.89489879, 162.92992838, 503.29199816, 483.86998168,  
        190.21357627, 320.16779413, 234.86640201, 329.93593349,  
        499.86680016, 312.36598191, 176.62649473, 347.92933836,  
        128.72291975, 121.32877046, 319.84723655, 266.87113957,  
        353.97563368, 333.89412921, 355.3973119 , 310.37518043,  
        210.16841178, 117.39194582, 169.21456175, 334.69348256,  
        365.6982803 , 220.81161853, 218.31586117,  99.22192384,  
        162.98307082, 299.55546367, 278.19363708, 286.69351682,  
        160.44166634, 508.02280332, 507.43020156, 182.74370679,  
        158.1028499 , 167.2093852 , 115.79663877, 193.70527458,  
        329.90516595, 191.52288727, 218.24633472, 115.33123245,  
        199.97832075,  98.38091523, 131.08061807, 123.56869251,  
        294.13102035, 136.07871873, 202.58855284, 346.99073276,  
        133.27149024, 206.87639605, 299.81229221, 232.79989585,  
        369.59455077, 140.60921235, 301.42777293, 168.49132713,  
        228.2135646 , 188.13948467, 181.21962139, 216.5724077 ,  
        508.83470261, 240.99140578, 394.08756449, 159.98535249,  
        113.92202733, 204.76344124, 132.18584505, 150.23000289,
```

```
170.43507112, 484.35797713, 205.4313322 , 340.87063536,
227.90768924, 216.04797057, 323.13225387, 396.92982461,
490.34416771, 158.87612585, 316.58500585, 146.37758745,
468.68248685, 135.64797219, 127.12382142, 247.51498683,
191.09763922, 128.98094863, 327.81528424, 343.50323443,
221.3701261 , 346.67226997, 322.73257353, 340.58454311,
318.48328648, 510.49850477, 151.80253551, 168.0221244 ,
211.21999526, 210.80783545, 116.61467072, 287.86836941,
202.99272334, 311.06854574, 305.4761803 , 491.98776778,
183.79289558, 180.40918037, 271.70590492, 293.79667101,
199.40422894, 352.78959461, 110.18958687, 129.41538464,
327.80359476, 336.5963633 , 492.94329592, 170.21949944,
499.66480094, 115.75608108, 204.66824552, 120.14812222,
285.1176868 , 496.04640213, 137.77661686, 96.39652271,
472.03171282, 518.38340016, 163.43339576, 136.71873802,
196.07655878, 219.88310127, 187.43303263, 304.91947564,
235.74746813, 158.22731696])
```

```
[55]: r2=r2_score(y_test, y_pred)
print("R2 for Random Forest Regressor Model is: ", r2)
```

R² for Random Forest Regressor Model is: 0.9998142074646916

```
[56]: mse=mean_squared_error(y_test, y_pred)
print("Mean Squared Error for Random Forest Regressor is : ", mse)
```

Mean Squared Error for Random Forest Regressor is : 2.3803557836653466

5 Using K-Nearest Neighbour Regressor model

```
[57]: from sklearn.neighbors import KNeighborsRegressor
```

```
[58]: #scaler=StandardScaler()
#X_train=scaler.fit_transform(X_train)
#x_test=scaler.fit_transform(X_test)
knn=KNeighborsRegressor(n_neighbors=4)
knn.fit(X_train, y_train)
```

```
[58]: KNeighborsRegressor(n_neighbors=4)
```

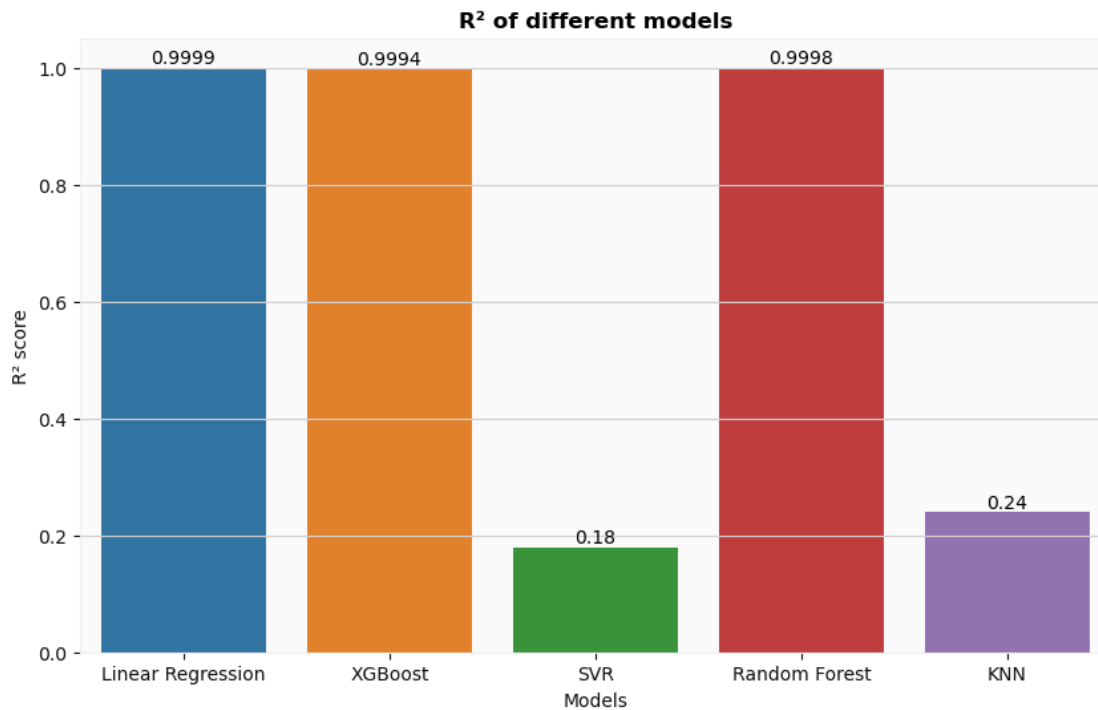
```
[59]: y_pred=knn.predict(X_test)
```

```
[60]: r2=r2_score(y_test, y_pred)
mse=mean_squared_error(y_test, y_pred)
print("R2 for KNN Regressor Model is : ", r2)
print("Mean Squared Error for KNN Regressor is : ", mse)
```

R^2 for KNN Regressor Model is : 0.24003349289411646
Mean Squared Error for KNN Regressor is : 9736.616530790738

```
[61]: models=["Linear Regression","XGBoost","SVR","Random Forest","KNN"]
r_square=[0.9999, 0.9994, 0.18, 0.9998, 0.24]

plt.figure(figsize=(10,6))
ax=sns.barplot(x=models, y=r_square)
for bar in ax.containers:
    ax.bar_label(bar)
plt.title("R2 of different models", fontweight='bold')
plt.xlabel("Models")
plt.ylabel("R2 score")
plt.show()
```



```
[62]: models=["Linear Regression","XGBoost","SVR","Random Forest","KNN"]
mean_square=[0.013, 7.21, 10472, 2.34, 9736]

plt.figure(figsize=(10,6))
ax=sns.barplot(x=models, y=mean_square)
for bar in ax.containers:
    ax.bar_label(bar)
plt.title("Mean Square Error of different models", fontweight='bold')
plt.xlabel("Models")
```

```
plt.ylabel("MSE score")
plt.show()
```



6 Conclusion

The analysis indicates that Linear Regression is most suitable model for the predication of the stock prices. Linear Regression Model has highest R^2 score .9999 as compared to Random Forest (.9998), XGBoost (.9994).

It also indicates that SVR and KNN are not the suitable models for such dataset.

Another metrics, Mean Square Error also indicates that Linear Regression has lowest MSE (0.013) making it best suitable model